

Module 6

You can download the R Markdown file (<https://gedeck.github.io/DS-6030/homework/Module-6.Rmd>) and use it to answer the following questions.

If not otherwise stated, use Tidyverse and Tidymodels for the assignments.

1. Predict out of state tuition (feature selection)

The data `College.csv` contains a number of variables for 777 different universities and colleges in the US. In this exercise, we will try to predict the `outstate` tuition fee using the other variables in the data set.

A. Data loading and preprocessing

```
#| message: FALSE
#| warning: FALSE
#| cache: TRUE

#Load packages
```

```
library(tidymodels)
```

```
## — Attaching packages —————— tidymodels 1.3.0 —
```

## ✓ broom	1.0.8	✓ recipes	1.3.0
## ✓ dials	1.4.0	✓ rsample	1.3.0
## ✓ dplyr	1.1.4	✓ tibble	3.2.1
## ✓ ggplot2	3.5.2	✓ tidyverse	1.3.1
## ✓ infer	1.0.8	✓ tune	1.3.0
## ✓ modeldata	1.4.0	✓ workflows	1.2.0
## ✓ parsnip	1.3.1	✓ workflowsets	1.1.0
## ✓ purrr	1.0.4	✓ yardstick	1.3.2

```
## — Conflicts —————— tidymodels_conflicts() —
## ✘ purrr::discard() masks scales::discard()
## ✘ dplyr::filter() masks stats::filter()
## ✘ dplyr::lag() masks stats::lag()
## ✘ recipes::step() masks stats::step()
```

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages ————— tidyverse 2.0.0 —
## ✓forcats 1.0.0 ✓readr 2.1.5
## ✓lubridate 1.9.4 ✓stringr 1.5.1
```

```
## — Conflicts ————— tidyverse_conflicts() —
## Xreadr::col_factor() masks scales::col_factor()
## Xpurrr::discard() masks scales::discard()
## Xdplyr::filter() masks stats::filter()
## Xstringr::fixed() masks recipes::fixed()
## Xdplyr::lag() masks stats::lag()
## Xreadr::spec() masks yardstick::spec()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(ISLR2)
library(mgcv)
```

```
## Loading required package: nlme
##
## Attaching package: 'nlme'
##
## The following object is masked from 'package:dplyr':
##   collapse
##
## This is mgcv 1.9-3. For overview type 'help("mgcv-package")'.
```

```
library(patchwork)
```

(1.1) Load the data from `ISLR2::College` and split into training and holdout sets using a 80/20 split. (1 point - coding)

```
set.seed(345)

#Load data

college <- ISLR2::College %>%
  as_tibble() %>%
  mutate(Private = as.factor(Private))

#Test/hold sets

set <- initial_split(college, prop = 0.8, strata = Outstate)
train <- training(set)
test <- testing(set)
```

B. Train linear regression and Lasso models

Use *tidymodels* to define a workflow and build a linear regression model to predict `Outstate` from all the other variables using L1 regularization (Lasso).

(1.2) For preprocessing, normalize all the numerical variables (`step_normalize(all_numeric_predictors())`) and convert the categorical / nominal variables to dummy variables (`step_dummy(all_nominal_predictors())`). (1 point - coding)

```
rec <- recipe(Outstate ~ ., data = train) %>%  
  step_normalize(all_numeric_predictors()) %>%  
  step_dummy(all_nominal_predictors())
```

(1.3) Train a normal linear regression model using the `lm` engine with the training set. Look at the *p*-values of the individual features. Which features are significant (use `extract_fit_engine` and `summary`). (1 point - coding/discussion)

```
lm_model <- linear_reg() %>%  
  set_engine("lm")  
  
lm_wf <- workflow() %>%  
  add_model(lm_model) %>%  
  add_recipe(rec)  
  
lm_fit <- lm_wf %>%  
  fit(data = train)  
  
#P-values  
  
lm_summary <- lm_fit %>%  
  extract_fit_engine() %>%  
  summary()  
  
lm_summary$coefficients
```

	##	Estimate	Std. Error	t value	Pr(> t)
##	(Intercept)	9003.09353	220.91238	40.7541368	3.041944e-175
##	Apps	-1252.48888	295.24569	-4.2421919	2.562231e-05
##	Accept	1944.32624	354.69494	5.4816859	6.203584e-08
##	Enroll	-115.77605	364.62948	-0.3175170	7.509615e-01
##	Top10perc	646.03278	222.00638	2.9099740	3.748143e-03
##	Top25perc	-201.42310	190.10963	-1.0595102	2.897924e-01
##	F.Undergrad	-769.86022	340.91939	-2.2581884	2.429113e-02
##	P.Undergrad	28.15111	107.84457	0.2610341	7.941555e-01
##	Room.Board	1090.19191	106.89445	10.1987698	1.228352e-22
##	Books	-109.92821	83.26308	-1.3202516	1.872528e-01
##	Personal	-139.71091	89.93623	-1.5534442	1.208426e-01
##	PhD	144.02247	156.52340	0.9201337	3.578714e-01
##	Terminal	445.91182	155.14606	2.8741422	4.194293e-03
##	S.F.Ratio	-272.95131	109.25008	-2.4984083	1.274077e-02
##	perc.alumni	526.14587	105.83532	4.9713637	8.680641e-07
##	Expend	922.44519	130.35158	7.0765937	4.124573e-12
##	Grad.Rate	422.52680	104.69126	4.0359319	6.140038e-05
##	Private_Yes	1968.53933	279.37023	7.0463461	5.043074e-12

Based on the p-values, it seems that the significant features are Grad.Rate, perc.alumni, Room.Board, Accept, Terminal, and Expend.

(1.4) Use `glmnet` and tune the L1 penalty parameter using 10-fold cross-validation. Tune the penalty over the range `penalty(c(-1, 2.5))` and check that the range is appropriate using `autoplot`. (2 points - coding)

```
lasso_model <- linear_reg(penalty = tune(), mixture = 1) %>%
  set_engine("glmnet")

lasso_wf <- workflow() %>%
  add_model(lasso_model) %>%
  add_recipe(rec)


#Grid

lasso_grid <- grid_regular(penalty(c(-1,2.5)), levels = 30)

#10-fold

set.seed(234)
lasso_res <- tune_grid(
  lasso_wf,
  resamples = vfold_cv(train, v = 10, strata = Outstate),
  grid = lasso_grid
)

autoplot(lasso_res)
```



(1.5) Determine the best penalty parameters using the lowest *RMSE* and the penalty obtained from the one-standard-error rule (`select_by_one_std_err` , see online material (<https://gedeck.github.io/DS-6030/book/model-tuning-examples.html#tune-one-standard-error>)). For both penalties, finalize the workflow and train models using the full training set. Report the coefficients of each model. Which variables are selected in each case? (2 points - coding/discussion)

```
#Lowest RMSE

best_rmse <- select_best(lasso_res, metric = "rmse")

#One standard error

one_se <- select_by_one_std_err(
  lasso_res, metric = "rmse", desc(penalty)
)

#Final workflows

final_rmse_wf <- finalize_workflow(lasso_wf, best_rmse)
final_one_se_wf <- finalize_workflow(lasso_wf, one_se)

#Models fit

final_rmse_fit <- fit(final_rmse_wf, data = train)
final_one_se_fit <- fit(final_one_se_wf, data = train)

#Coefficients of models

coef_rmse <- tidy(final_rmse_fit)
coef_one_se <- tidy(final_one_se_fit)

coef_rmse
```

```
## # A tibble: 18 × 3
##   term      estimate penalty
##   <chr>     <dbl>    <dbl>
## 1 (Intercept) 8947.     14.9
## 2 Apps        -759.     14.9
## 3 Accept       1313.    14.9
## 4 Enroll        0       14.9
## 5 Top10perc    364.    14.9
## 6 Top25perc   -0.0197  14.9
## 7 F.Undergrad  -679.    14.9
## 8 P.Undergrad   0       14.9
## 9 Room.Board    1089.   14.9
## 10 Books       -96.9    14.9
## 11 Personal     -140.   14.9
## 12 PhD          167.    14.9
## 13 Terminal     425.    14.9
## 14 S.F.Ratio   -273.   14.9
## 15 perc.alumni  517.    14.9
## 16 Expend       927.    14.9
## 17 Grad.Rate    414.    14.9
## 18 Private_Yes 2044.   14.9
```

```
coef_one_se
```

```
## # A tibble: 18 × 3
##   term      estimate penalty
##   <chr>     <dbl>    <dbl>
## 1 (Intercept) 9011.    240.
## 2 Apps          0       240.
## 3 Accept         0       240.
## 4 Enroll         0       240.
## 5 Top10perc    140.    240.
## 6 Top25perc     0       240.
## 7 F.Undergrad    0       240.
## 8 P.Undergrad    0       240.
## 9 Room.Board    1092.   240.
## 10 Books         0       240.
## 11 Personal     -49.7   240.
## 12 PhD           112.   240.
## 13 Terminal      310.   240.
## 14 S.F.Ratio    -225.   240.
## 15 perc.alumni   521.   240.
## 16 Expend        923.   240.
## 17 Grad.Rate     431.   240.
## 18 Private_Yes   1957.  240.
```

It seems that the One standard error model selects less variables than the RMSE model, thus making the model more regularized. All the variables were selected for the One standard error model except Apps, Accept, Enroll, Top25Perc, F.Undergrad, P.Undergrad, and Books.

(1.6) Use the model from (1.3) and the two models from (1.5) to predict the `Outstate` variable on the training and test set. Report the RMSE and R^2 of each model on the training and test set. (1 point - coding/discussion)

```
get_metrics <- function(model, train, test) {
  bind_rows(
    metrics(model %>% augment(train), truth = Outstate, estimate = .pred) %>%
      mutate(dataset = "train"),
    metrics(model %>% augment(test), truth = Outstate, estimate = .pred) %>%
      mutate(dataset = "test")
  )
}

#Metrics

lm_metrics <- get_metrics(lm_fit, train, test) %>% mutate(model = "Linear")
rmse_metrics <- get_metrics(final_rmse_fit, train, test) %>% mutate(model = "Lasso RSME")
one_se_metrics <- get_metrics(final_one_se_fit, train, test) %>% mutate(model = "Lasso 1SE")

all_metrics <- bind_rows(lm_metrics, rmse_metrics, one_se_metrics)

all_metrics %>%
  pivot_wider(names_from = .metric, values_from = .estimate)
```

```
## # A tibble: 6 × 6
##   .estimator dataset model      rmse    rsq    mae
##   <chr>     <chr>  <chr>      <dbl> <dbl> <dbl>
## 1 standard  train   Linear     1934.  0.773 1520.
## 2 standard  test    Linear     1934.  0.773 1520.
## 3 standard  train   Lasso RSME 1941.  0.772 1524.
## 4 standard  test    Lasso RSME 1941.  0.772 1524.
## 5 standard  train   Lasso 1SE  2037.  0.756 1598.
## 6 standard  test    Lasso 1SE  2037.  0.756 1598.
```

Based on the values of the RSME and R² for each metric, it seems that linear regression and RSME perform well, but there could be some possible overfitting. The One standard error shows good performance as well, but has less possibility of overfitting.

2. Predict out of state tuition (GAM model)

C. GAM model

Using the significant features from (1.3), build a generalized additive model (GAM) to predict `Outstate`. (see Generalized additive models (GAM) (https://gedeck.github.io/DS-6030/book/deep-dive-gen_additive_mod.html) for how to build GAM models in `tidymodels`)

(2.1) Define a model formula setting all numerical variables as splines and all categorical variables as factors
`(Outstate ~ Private + s(Apps) + ...)` (1 point - coding)

```
gam_formula <- Outstate ~ Private + s(Grad.Rate) + s(perc.alumni) + s(Room.Board) + s(Expend) +
  s(Terminal) + s(Accept)
```

(2.2) Define the `gen_additive_mod` model using `mgcv` as the engine and fit the model using the training data (2 points - coding)

```
gam_model <- gen_additive_mod() %>%
  set_engine("mgcv") %>%
  set_mode("regression")

#Model fit

gam_fit <- fit(gam_model, formula = gam_formula, data = train)
```

(2.3) Report the RMSE and R^2 of the model on the training and test set and compare with the result from (1.6) (1 point - discussion)

```
gam_metrics <- get_metrics(gam_fit, train, test) %>%
  mutate(model = "GAM")

gam_metrics %>%
  pivot_wider(names_from = .metric, values_from = .estimate)
```


```
## # A tibble: 2 × 6
##   .estimator dataset model  rmse    rsq    mae
##   <chr>      <chr>   <chr> <dbl> <dbl> <dbl>
## 1 standard   train    GAM   1782.  0.807 1367.
## 2 standard   test     GAM   1782.  0.807 1367.
```



Based on the metric values from GAM, it looks like GAM performed better than the Linear and Lasso models. This is most likely due to its ability to capture non-linear relationships.



(2.4) Use the `plot` function with the fitted model (use `extract_fit_engine` to get the actual `mgcv` model for plotting; set `scale=0` to get individual y -scales). Describe your observations. (1 point - coding/discussion)


```
gam_engine <- gam_fit %>%
  extract_fit_engine()

plot(gam_engine, scale = 0)
```









The plots show smooth, non-linear relationships between predictors Accept, Expend, and Grad.Rate.

(2.5) Use the `summary` function to get information about the model. Based on the reported significance levels, could you simplify the model further? (1 point - coding/discussion)

```
summary(gam_engine)
```

```

## 
## Family: gaussian
## Link function: identity
##
## Formula:
## Outstate ~ Private + s(Grad.Rate) + s(perc.alumni) + s(Room.Board) +
##   s(Expend) + s(Terminal) + s(Accept)
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 8646.5     205.8  42.006 <2e-16 ***
## PrivateYes  2451.2     260.6   9.407 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df      F p-value
## s(Grad.Rate) 3.814 4.806 5.932 5.37e-05 ***
## s(perc.alumni) 1.396 1.701 13.743 1.00e-05 ***
## s(Room.Board) 1.225 1.420 52.151 < 2e-16 ***
## s(Expend)    6.760 7.851 33.636 < 2e-16 ***
## s(Terminal)   1.000 1.000  3.806  0.0515 .
## s(Accept)     6.461 7.541  2.540  0.0110 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) = 0.8 Deviance explained = 80.7%
## GCV = 3.4193e+06 Scale est. = 3.2944e+06 n = 620

```

Based on the p-values, it seems as though Terminal could be taken out to simplify the model.

(2.6) Simplify the model by removing the non-significant variables and re-fit the model. Report the RMSE and R^2 of the model on the training and test set. (1 point - coding/discussion)

```

gam_formula_simple <- Outstate ~ Private + s(Grad.Rate) + s(perc.alumni) + s(Room.Board) + s(Expend) + s(Accept)

gam_model_simple <- gen_additive_mod() %>%
  set_engine("mgcv") %>%
  set_mode("regression")

#Model fit

gam_fit_simple <- fit(gam_model_simple, formula = gam_formula_simple, data = train)

# Check performance
gam_simple_metrics <- get_metrics(gam_fit_simple, train, test) %>%
  mutate(model = "GAM Simplified")

gam_simple_metrics %>%
  pivot_wider(names_from = .metric, values_from = .estimate)

```

```
## # A tibble: 2 × 6
##   .estimator dataset model      rmse    rsq    mae
##   <chr>     <chr>   <chr>      <dbl> <dbl> <dbl>
## 1 standard  train   GAM Simplified 1786. 0.807 1374.
## 2 standard  test    GAM Simplified 1786. 0.807 1374.
```

The RMSE for the training and test set is 1785.732, and the R² for the training and test set is 0.8066 for both sets.

D. Comparison

(2.7) Compare the results from the three models from (B), the full GAM model from (2.2) and the reduced model from (2.6)? (2 points - discussion)

```
all_model_matrices <- bind_rows(
  all_metrics,
  gam_metrics,
  gam_simple_metrics
)

all_model_matrices %>%
  pivot_wider(names_from = .metric, values_from = .estimate)
```

```
## # A tibble: 10 × 6
##   .estimator dataset model      rmse    rsq    mae
##   <chr>     <chr>   <chr>      <dbl> <dbl> <dbl>
## 1 standard  train   Linear    1934. 0.773 1520.
## 2 standard  test    Linear    1934. 0.773 1520.
## 3 standard  train   Lasso RSME 1941. 0.772 1524.
## 4 standard  test    Lasso RSME 1941. 0.772 1524.
## 5 standard  train   Lasso 1SE  2037. 0.756 1598.
## 6 standard  test    Lasso 1SE  2037. 0.756 1598.
## 7 standard  train   GAM      1782. 0.807 1367.
## 8 standard  test    GAM      1782. 0.807 1367.
## 9 standard  train   GAM Simplified 1786. 0.807 1374.
## 10 standard test   GAM Simplified 1786. 0.807 1374.
```

From all of the metrics, it seems that the full GAM model has the best performance. It has the lowest RMSE, as well as the highest R². The simplified GAM also seems to have performed better than the Linear and One standard error models, but doesn't perform as well as the full GAM model.